

Scanning/Lexical Analysis

CIS*4650 (Winter 2020)

Review

- **Compiler:** translates a source program into a target machine code
 - **Front-End Analysis:** language-specific, focused on analysis, and systematic solutions with tools
 - **Back-End Synthesis:** machine-specific, focused on synthesis, and ad hoc solutions through coding
- **Front-End Analysis:**
 - Scanning/Lexical Analysis: break an input stream into a token sequence
 - Parsing/Syntactic Analysis: parse the phrase structure of a program
 - Semantic Analysis: calculate the meaning and generate the intermediate code of a program

What is a Scanner?

- Split input characters into a sequence of tokens:



What is a Token?

- A minimal sequence of characters that represents a unit of information
- Each language specifies a finite set of token types

Type	Examples
ID	foo n14 last
INTEGER	73 0 515 082
REAL	66.1 .5 10. 1e67 5.5e-10
IF	if
COMMA	,
NOTEQ	!=
LPAREN	(
RPAREN)

Non-token Examples

- Some character sequences are not tokens:

comment

```
/* try again */
```

preprocessor directive

```
#include <stdio.h>
```

preprocessor directive

```
#define NUMS 5 , 6
```

macro

```
NUMS
```

blanks, tabs, newlines

Example Input/Output

● Input program:

```
/* find a zero */  
float match0(char *s)  
{  
    if( !strncmp(s, "0.0", 3) )  
        return(0.0);  
}
```

● Output tokens:

```
FLOAT ID(match0) LPAREN CHAR STAR ID(s) RPAREN LBRACE  
IF LPAREN BANG ID(strncmp) LPAREN ID(s) COMMA  
STRING(0.0) COMMA INTEGER(3) RPAREN RPAREN RETURN  
REAL(0.0) SEMI RBRACE EOF
```

Specifying Tokens

- Token structures can be complex, and English descriptions can be tedious, imprecise, and incomplete
 - E.g., identifiers and real numbers
- Need a formal system to specify them without ambiguity:
 - Allow review of design and validation of implementation
- Regular expressions:
 - succinct, precise
 - capable of representing infinite sets of strings

English Rules for Identifiers

An identifier is a sequence of letters and digits; the first character must be a letter.

The underscore `_` counts as a letter.

Upper- and lower-case letters are different.

If the input stream has been parsed into tokens up to a given character, the next token is taken to include the longest string of characters that could possibly constitute a token.

Blanks, tabs, newlines, and comments are ignored except as they serve to separate tokens.

Some white space is required to separate otherwise adjacent identifiers, keywords, and constants.

Regular Expressions

● Basic Regular Expressions:

➤ Given any character a from an alphabet Σ , a itself denotes a basic regular expression and the language it recognizes is $L(a) = \{a\}$.

➤ Empty string: ε (also λ)

$$L(\varepsilon) = \{\varepsilon\}$$

➤ Empty set: Φ

$$L(\Phi) = \{\}$$

➤ What is the difference between ε and Φ ?

Regular Expressions

● Basic Operations:

➤ concatenation: ab (also $a \cdot b$)

$$L(ab) = L(a) \times L(b) = \{ab\}$$

➤ choice/alternation: $a \mid b$

$$L(a \mid b) = L(a) \cup L(b) = \{a, b\}$$

➤ Kleene closure: a^*

$$L(a^*) = L(\epsilon) \cup L(a) \cup L(aa) \cup L(aaa) \dots = \{\epsilon, a, aa, aaa, \dots\}$$

➤ Combinations:

$$(a \mid b)a = \{aa, ba\}$$

$$((a \mid b)a)^* = \{\epsilon, aa, ba, aaaa, aaba, baaa, baba, aaaaaa, \dots\}$$

Regular Expressions

● Example 1:

- Regular expression: $b^*(abb^*)^*(a | \epsilon)$
- Language: strings of a's and b's with no consecutive a's

● Example 2:

- Language: strings of a's and b's containing consecutive a's
- Regular expression: $(a | b)^*aa(a | b)^*$

● Precedence:

- Highest to lowest: closure, concatenation, and choice.
- Without precedence: $a | b^* = (a | b)^*$ or $a | b^* = a | (b^*)?$
- With precedence: $a | b^* = a | (b^*)$

Extended Notations

- One or more repetitions: $a^+ = L(aa^*) = \{a, aa, aaa, \dots\}$
- Any character: $.$ = {any character in the alphabet}
- Range of characters:
 - $[a-z]$ = {all lowercase letters}
 - $[a-zA-Z]$ = {all lowercase and uppercase letters}
- Any characters not in the alphabet:
 - $\sim(a|b|c)$ = {any character that is not a or b or c}
- Optional expression: $a? = a|\epsilon$
- Strings: “ $a.+$ ” = {the string itself}

Regular Expressions for Tokens

Some token categories:

if	IF-TOKEN
[a-z][a-z0-9]*	ID-TOKEN
[0-9]+	NUM-TOKEN
([0-9]+ "." [0-9]*) ([0-9]* "." [0-9]+)	REAL-TOKEN
("--" [a-z]* "\n") (" " "\n" "\t")+	WHITE-SPACE
.	ERROR

Disambiguation Rules:

- Longest match: e.g., if8 is taken as an identifier
- Rule-Order Priority: e.g., if is taken as a reserved word.

Requirements for a Scanner

- The set of patterns should form a partition of all possible token classes: mutually exclusive and exhaustive



- No two classes have any intersections
- The union of all classes fill up the whole universe

- Implication: any stream of characters can be tokenized, and each token only belongs to one class
 - A filler pattern at the end is often needed for certain scanners

Deterministic Finite Automata

- Regular expressions do not provide a control mechanism for implementation
- Deterministic Finite Automaton (DFA): a simple machine that recognizes strings of regular sets
 - A finite set of states S
 - A finite alphabet Σ
 - A transition function $T: S \times \Sigma \rightarrow S$
 - A specific start state s_0 in S
 - A set of accepting or final states: $F \subseteq S$
- The set of strings accepted by a DFA defines a language

Graphical Representation

$$S = \{s0, s1\}$$

$$\Sigma = \{a, b\}$$

$$F = \{s1\}$$


$$T(s0, a) = s1$$

$$T(s0, b) = s0$$


$$T(s1, a) = s1$$

$$T(s1, b) = s0$$


● Given the following DFA:




DFAs for Tokens




IF




ID




NUM



REAL

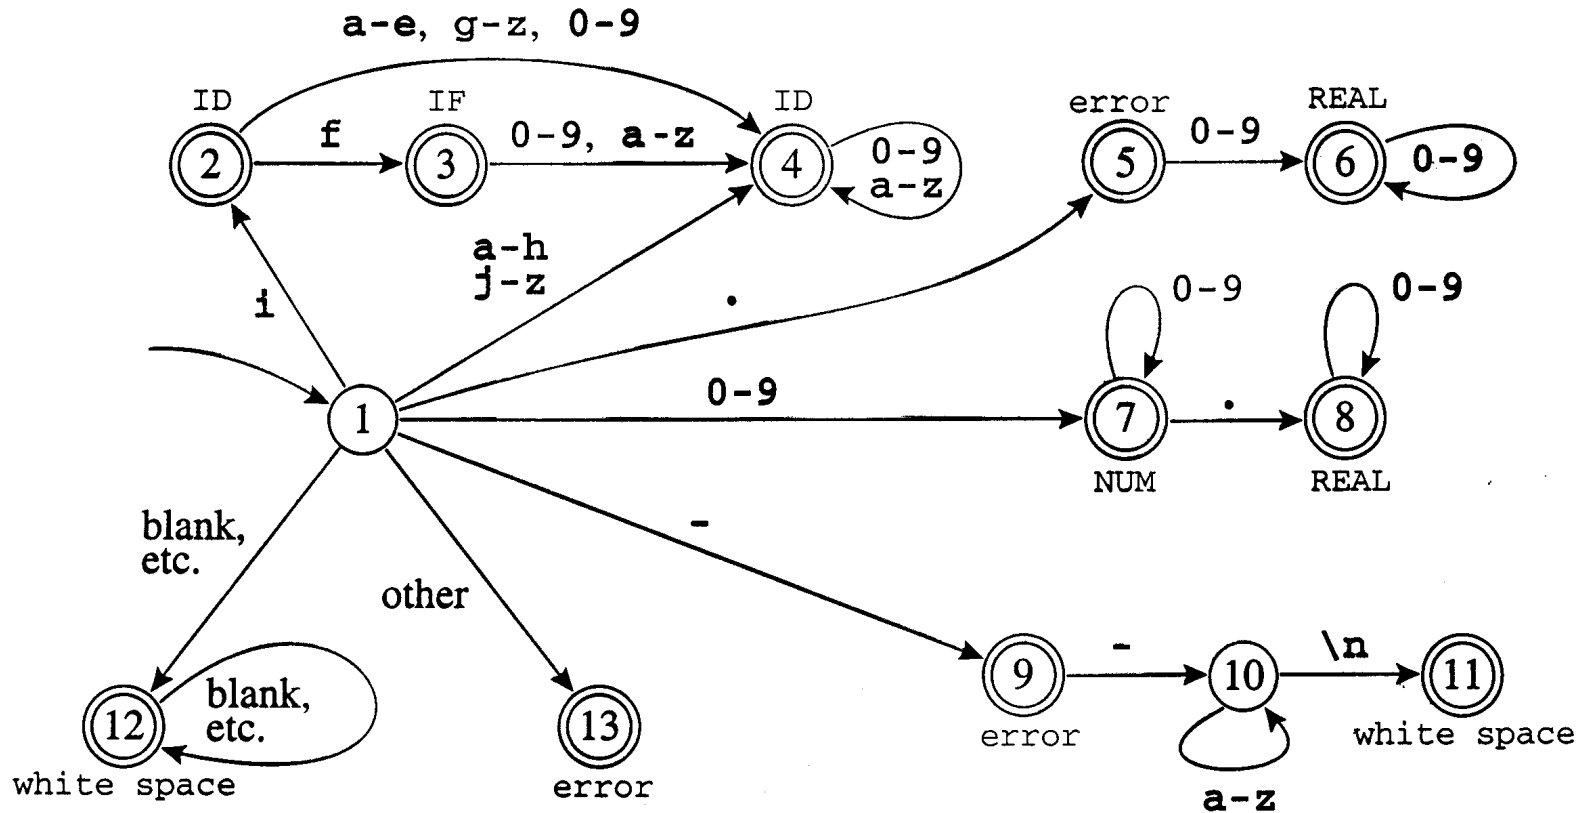


white space



error


Combined DFA



Non-Deterministic Finite Automata



● An NFA is different from a DFA:

- More than one edge from a given state can be labeled with the same input symbol.
- Edges can be labeled with ϵ (transition without consuming input).




Regular Expressions to NFAs

- DFAs and NFAs have the same expressive power (i.e., accepting only regular sets).
- Why bother with NFAs?
 - Easier to convert regular expressions to NFAs
- Thompson's Construction:




Regular Expressions and NFAs

Concatenation: ab




Choice: $a \mid b$

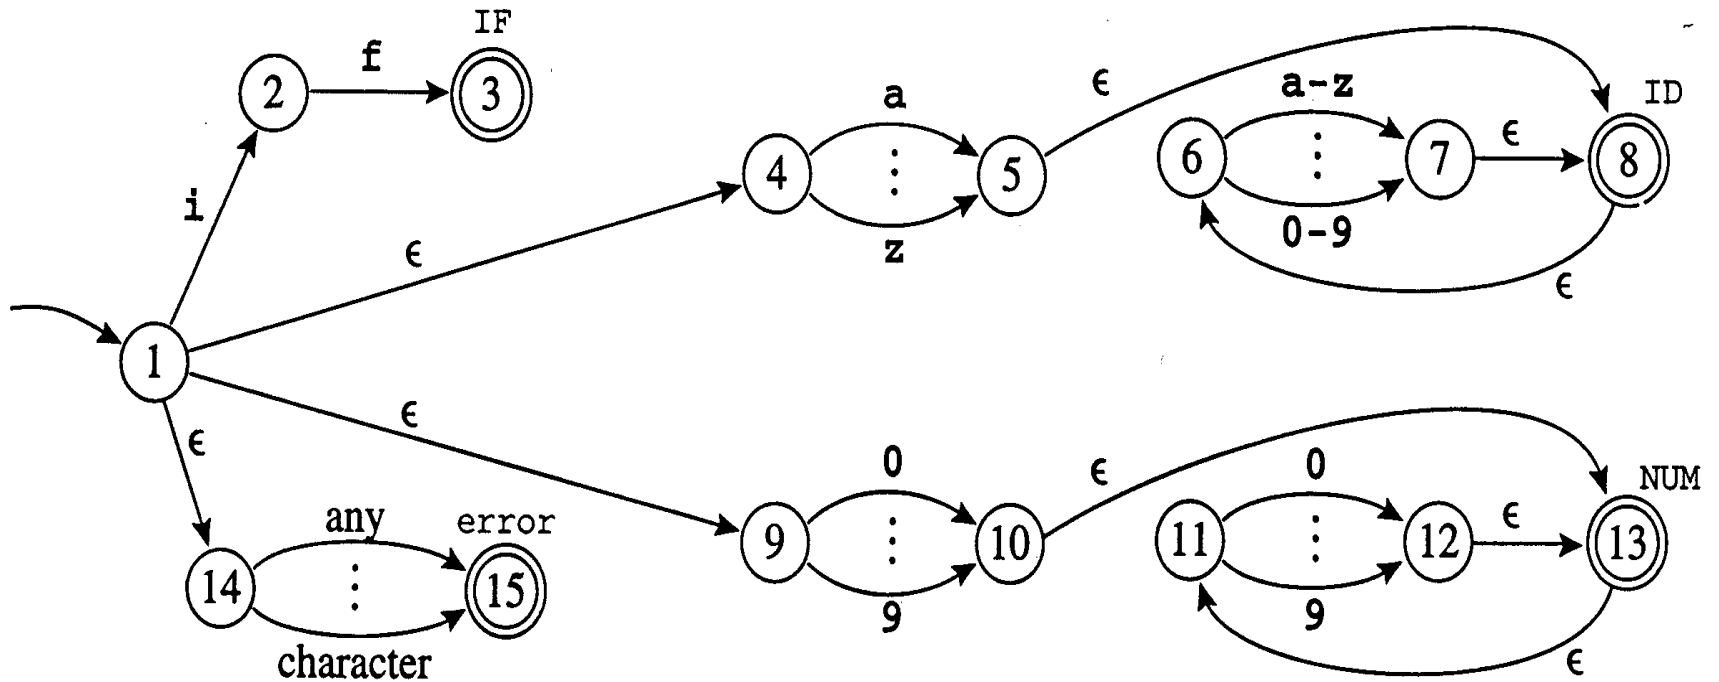


Regular Expressions and NFAs

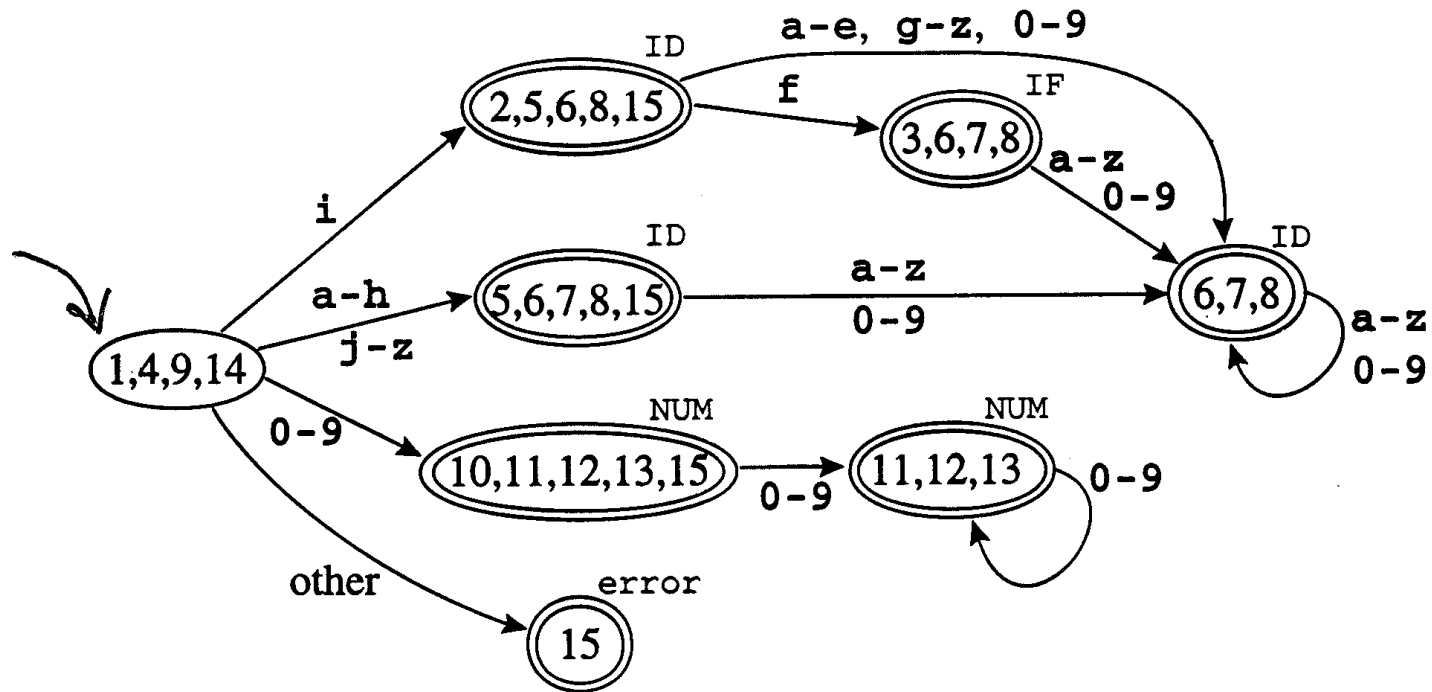
● Kleene Closure: a^*



NFA for Tokens



Converted DFA




Desirable and Optimal Solution

- Formulate all the regular expressions
- Convert regular expressions into a combined NFA
- Convert the NFA to a DFA
- Minimize/optimize the DFA for the computation

DFA: Case-Based Implementation

```
state := 1; ch := next input char;  
while( state != 0 and ch != EOF ) do  
  case state of  
    1: case ch of  
      letter: state := 2; ch := next input char;  
      else state := 0;  
    2: case ch of  
      letter, digit: state := 2; ch := next input char;  
      else state := 0;  
  end while;  
if( ch == EOF and state is final )  
  accept;  
else  
  report error;
```



DFA: Table-Based Implementation

```
state := 1; ch := next input char;  
while( state != 0 and ch != EOF ) do  
    state := T[state, ch];  
    if( state != 0 )  
        ch := next input char;  
end while;  
if( ch == EOF and state is final )  
    accept;  
else  
    report error;
```

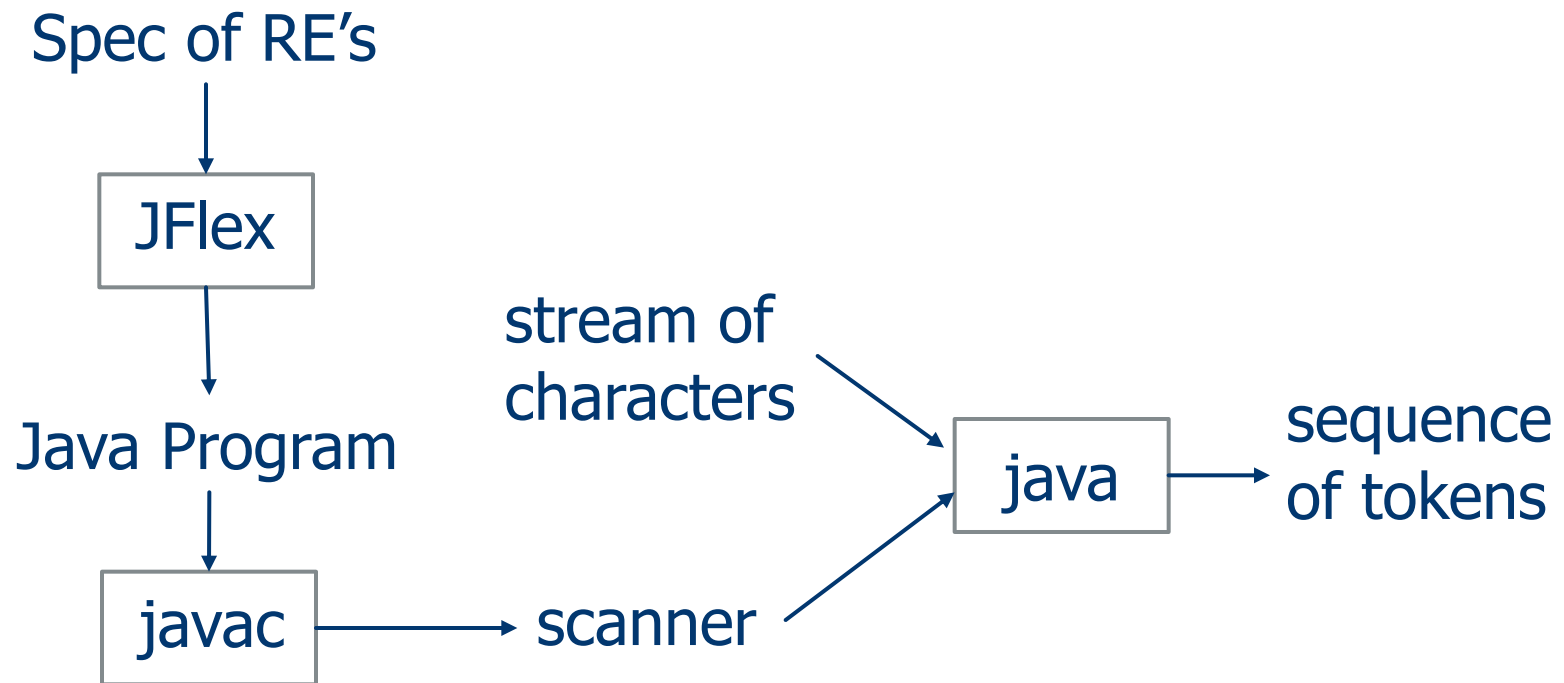
input state	letter	digit	other
1	2		
2	2	2	

Scanner Generator

- Writing scanners is a common requirement
 - Parsing is a ubiquitous task
- Process is repetitive, resulting in similar code structure
- Process is not hard to automate
- Scanner generators receive a specification file
 - definitions of tokens to be scanned
 - non-procedural programming: focus on what, not how
 - e.g., Lex, Flex, and JFlex.

Scanner Generator

- Automatically convert a specification file into a program that implements a scanner



Sample Jflex Specification (1 / 2)

```
/* JFlex specification for LISTL language */
```

```
%%
```

```
%class Lexer
```

```
%type Token
```

```
%line
```

```
%column
```

```
%eofval{
```

```
    return null;
```

```
%eofval}
```

```
digit = [0-9]
```

```
number = {digit}+
```

```
newline = \n|\r|\r\n
```

```
whitespace = [ \t]+
```

```
sign = ("+"|"-"")?
```

Sample Jflex Specification (2/2)

%%


```
{sign}{number}""? {  
    return new Token(Token.INTEGER, yytext(), yyline, yycolumn); }  
{sign}{digit}*""{number}(e{sign}{number})? {  
    return new Token(Token.FLOAT, yytext(), yyline, yycolumn); }  
{sign}{number}""/{number} {  
    return new Token(Token.RATIO, yytext(), yyline, yycolumn); }  
  
"(" { return new Token(Token.LPAREN, yytext(), yyline, yycolumn); }  
")" { return new Token(Token.RPAREN, yytext(), yyline, yycolumn); }  
car { return new Token(Token.CAR, yytext(), yyline, yycolumn); }  
cdr { return new Token(Token.CDR, yytext(), yyline, yycolumn); }  
  
{newline} { /* skip newline */ }  
{whitespace} { /* skip whitespaces */ }  
[^ \t\n()]+ { return new Token(Token.SYMBOL, yytext(), yyline, yycolumn); }
```

Token Class

```
class Token {  
    public static final int INTEGER = 0;  
    public static final int FLOAT = 1;  
    public static final int RATIO = 2;  
    public static final int LPAREN = 3;  
    public static final int RPAREN = 4;  
    public static final int CAR = 5;  
    public static final int CDR = 6;  
    public static final int SYMBOL = 7;  
  
    public int index;  
    public String value;  
    public int line;  
    public int column;  
  
    Token( int index, String value, int line, int column ) {  
        this.index = index; this.value = value; this.line = line; this.column = column;  
    }  
}
```


Macro States in Scanner Tools

● Using Macro States:



Macro States in JFlex

```
/* the usual preamble ... */
```

```
%%
```

```
/* some definitions */
```

```
%state COMMENT
```

```
/* other definitions */
```

```
%%
```

```
/* regular expressions and actions */
```

```
<YYINITIAL>if { return symbol(sym.IF); }
```

```
<YYINITIAL>[a-z][a-z0-9]* { return symbol(sym.ID, yytext()); }
```

```
<YYINITIAL>"/*" { yybegin(COMMENT); }
```

```
<YYINITIAL>[ \t\n]* { /* whitespaces */ }
```

```
<YYINITIAL>.{ return symbol(sym.ERROR, yytext()); }
```

```
<COMMENT>"*/" { yybegin(YYINITIAL); }
```

```
<COMMENT>.\n { /* skip comments */ }
```